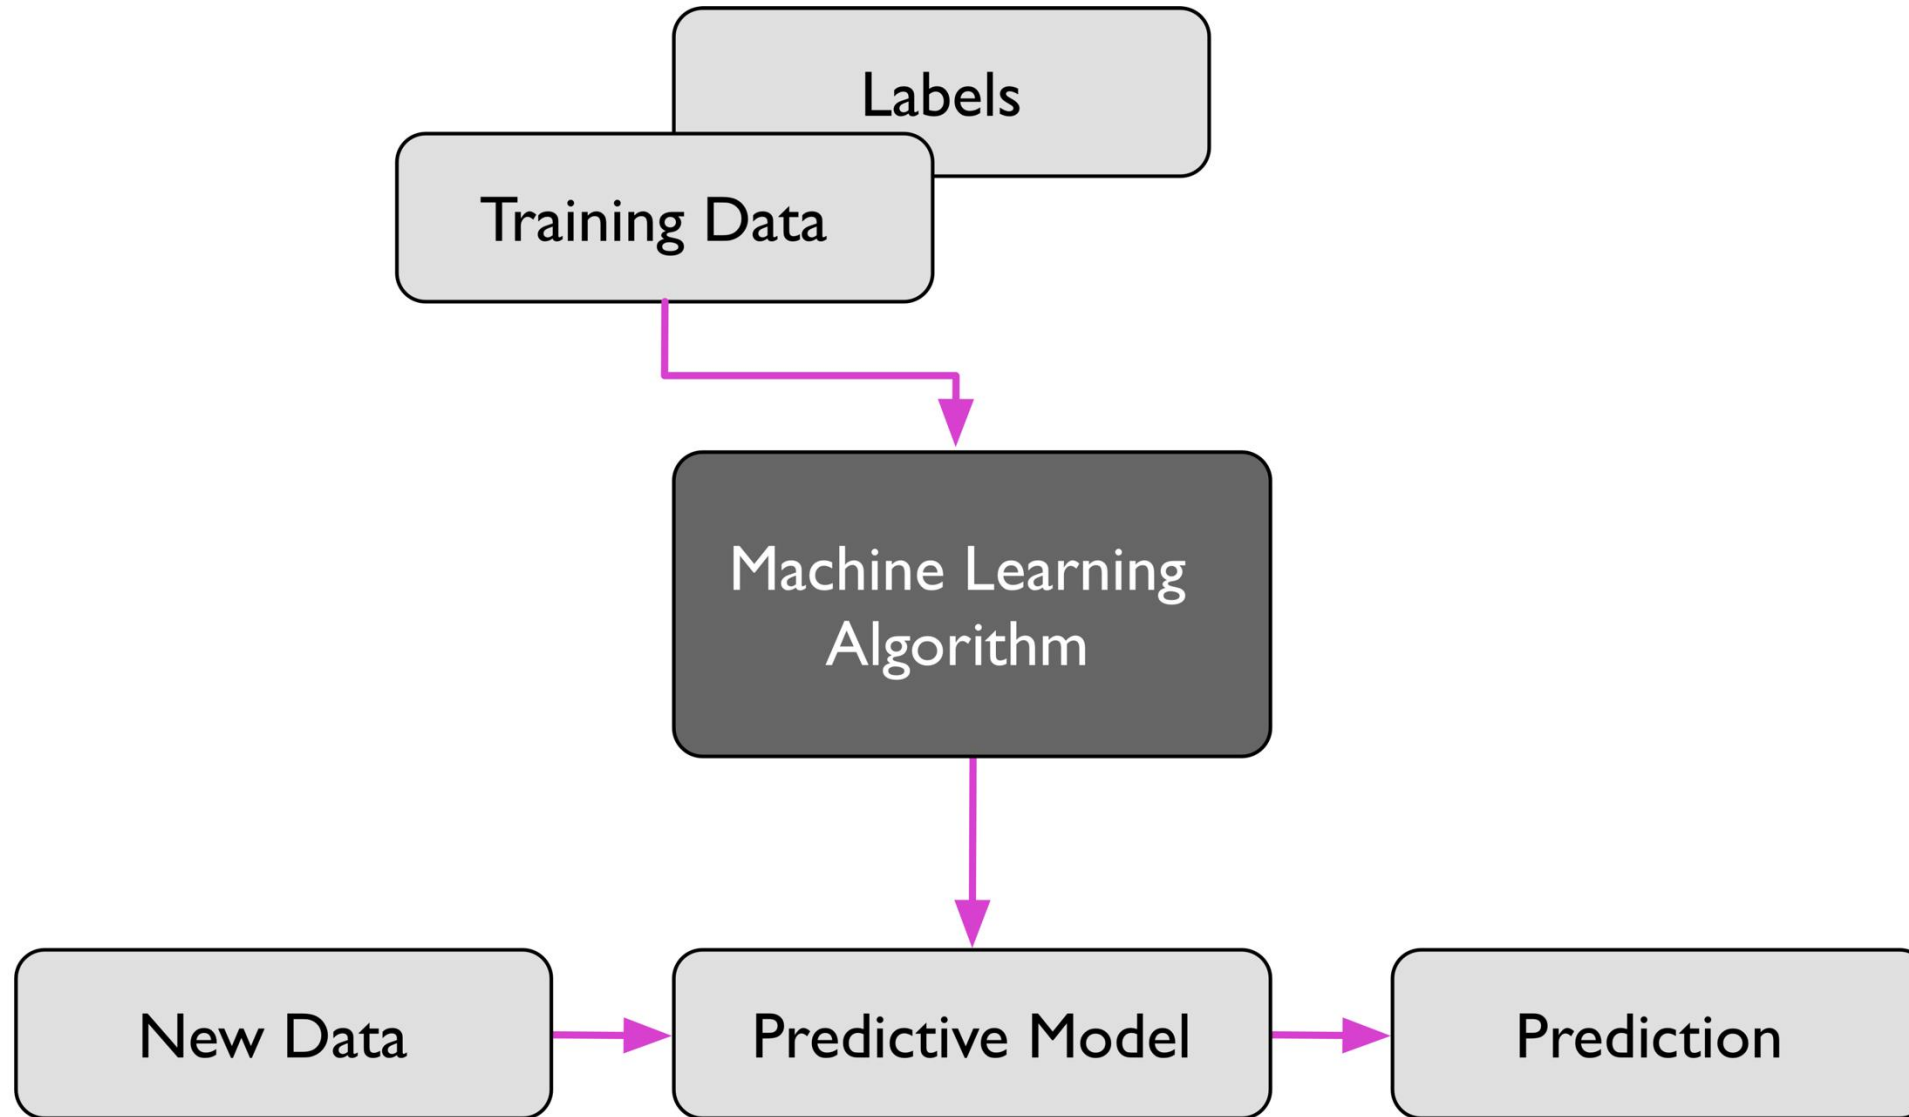


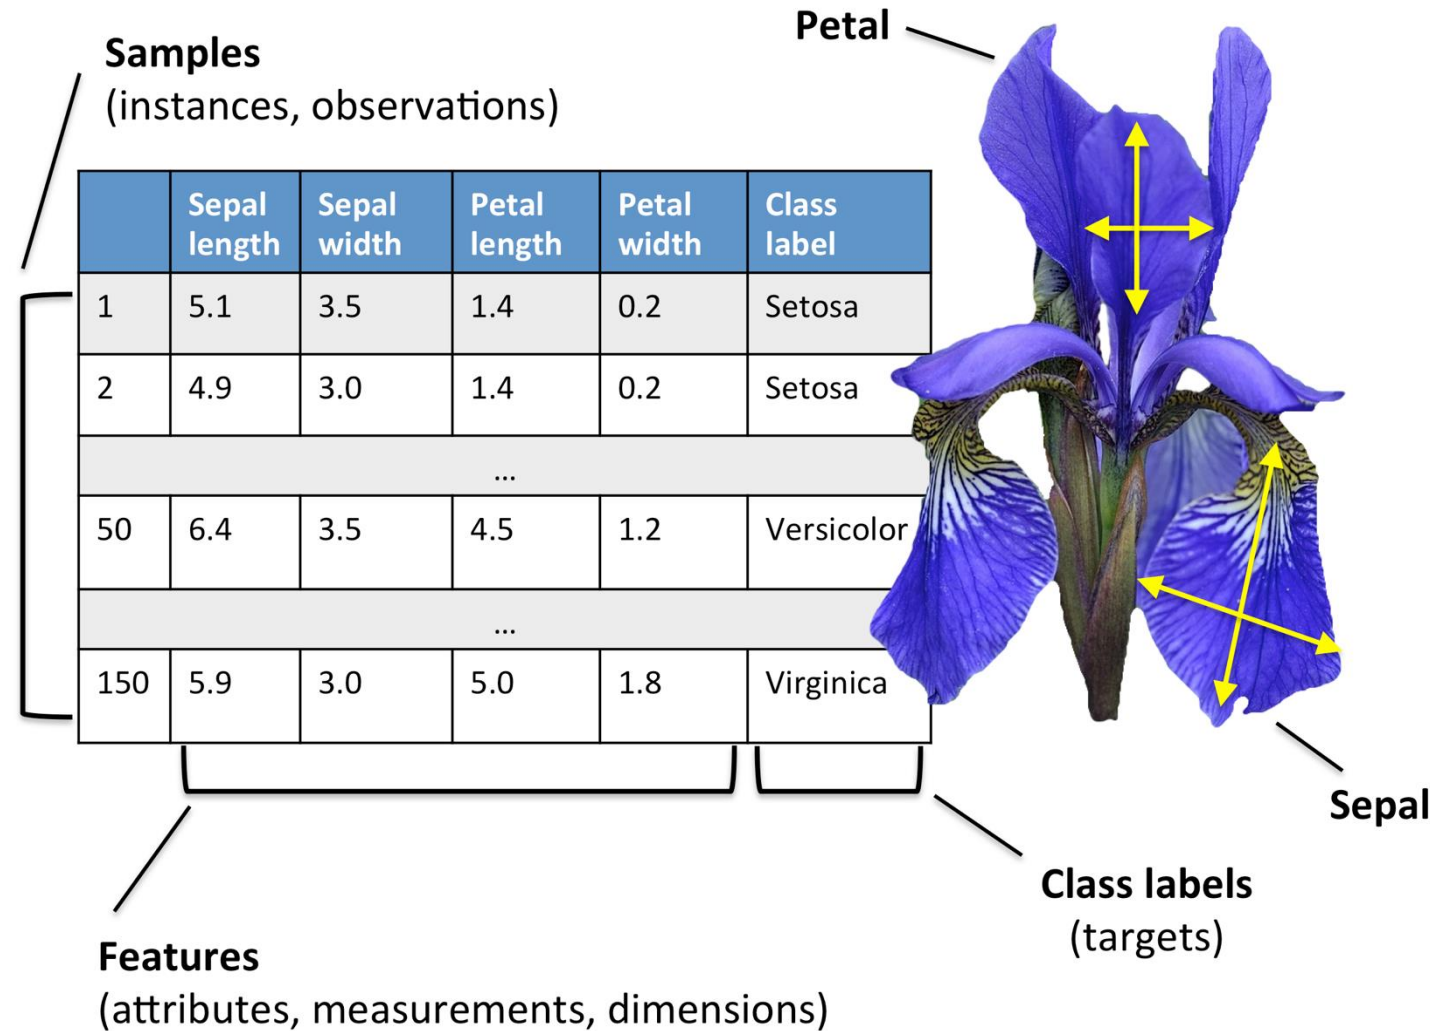
# Three core types of ML problems

---



# Generic supervised learning problem



# Generic supervised learning problem



# Classification vs regression



# Why ML models fail

---

# How much data is enough?


---

# What about features?

---

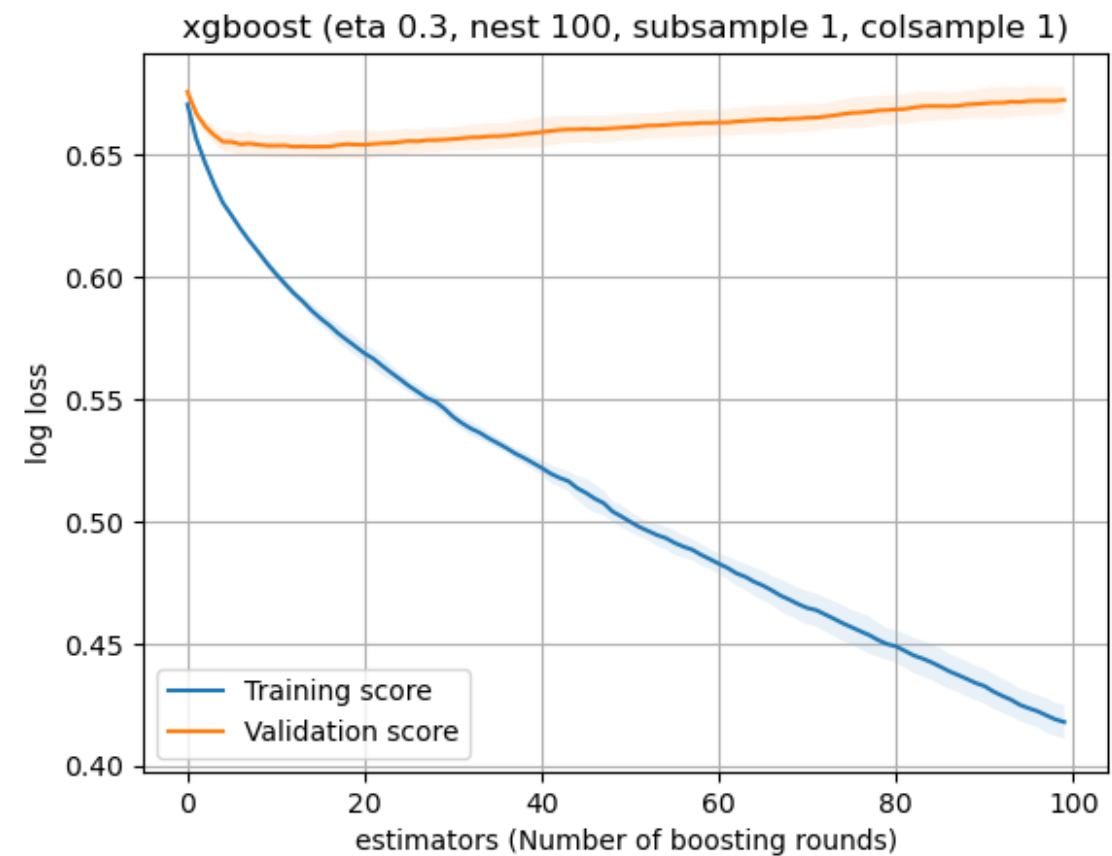

What are attributes of **good/bad** features?

# Validation is essential!





# Validation is essential!



# How to approach an ML problem

---

What is the right property to predict?

What is the right way to assess performance?

Which data points are compatible?

Which features might be informative?

...